

---

# Contents

---

- **LVM (Logical Volume Manager)**
  - **What is LVM (Logical Volume Manager) ?**
  - **Anatomy of LVM**
  - **Mapping modes (linear/striped)**
  - **Snapshots**
  - **Boot time scripts**
  - **LVM Files**
  - **LVM Commands**
  - **Examples**
  - **Kernel configuration**


---

# What is LVM (Logical Volume Manager) ?

---

- **LVM** provides a higher-level view of the disk storage on a computer system than the traditional view of disks and partitions.
- Storage volumes created under the control of the LVM can be resized and moved around almost at will.

# Anatomy of LVM



**VG (Volume Group):** gathers together LVs and PVs into one administrative unit.

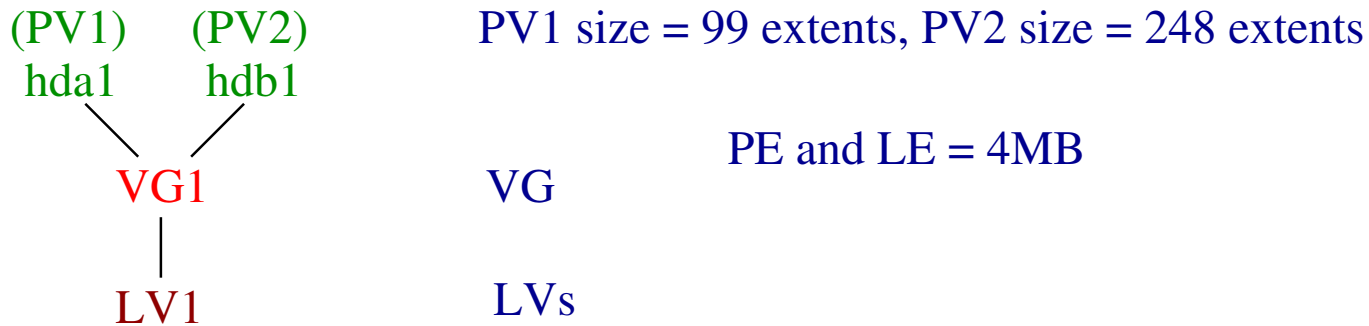
**PV (Physical Volume):** a PV is typically a hard disk (partition), raid device, etc...

**LV (Logical Volume):** the equivalent of a disk partition in a non-LVM system.

**PE (Physical Extent):** each PV is divided in chunks of data, known as PEs. These extents have the same size as the LEs.

**LE (Logical Extent):** each LV is split into chunks of data, known as LEs.

# Mapping modes (linear/striped) 1/2



- We now can create a LV (LV1).
- LV1 can be any size between 1 and 347 (99+248) extents.
- When LV1 is created, a mapping is defined between LEs and PEs (eg. LE[i] could map onto PE[i] of PV1 and PE[i] of PV2).

---

# Mapping modes (linear/striped) 2/2

---

**Strategies** for mapping LEs onto PEs:

1. **Linear** ( $\simeq$  RAID Linear) mapping will assign a range of PEs to an area of an LV in order (eg., LE[1 - 99] map onto PV1 and LE[100 - 248] map onto PV2).

2. **Striped** ( $\simeq$  RAID0):

LE[1]  $\rightarrow$  PV1[1],

LE[2]  $\rightarrow$  PV2[1],

LE[3]  $\rightarrow$  PV1[2],

and so on.

---

# Snapshots

---

- **Snapshots** allows the administrator to create a new block device which is an exact copy of a LV, frozen at some point in time.
- Used when for instance, we need to perform a backup on the LV, but you don't want to halt a live system that is changing the data.

---

# Boot time scripts 1/3

---

- The **startup** of LVM requires just the following two commands:

```
# vgscan
```

```
vgscan – reading all physical volumes (this may take a while...)
```

```
vgscan – "/etc/lvmtab" and "/etc/lvmtab.d" successfully created
```

```
vgscan – WARNING: This program does not do a VGDA backup of your volume group
```

```
# vgchange -ay
```

- And the **shutdown** only one:

```
# vgchange -an
```

Follow the next instructions depending on the Linux distribution you are running

# Boot time scripts - Debian (lvm tool package)

## 2/3

---

Create a startup script in `/etc/init.d/lvm` containing the following:

```
#!/bin/sh
case "$1" in
    start) /sbin/vgscan
           /sbin/vgchange -ay ;;
    stop) /sbin/vgchange -an ;;
    restart|force-reload)
           /sbin/vgchange -an && /sbin/vgscan \\\
           && /sbin/vgchange -ay ;;
esac
exit 0
```



---

# Boot time scripts - Redhat 3/3

---

## 1. Modify `/etc/rc.d/rc.sysinit` to load the LVM:

```
# LVM initialization
if [ -e /proc/lvm -a -x /sbin/vgchange -a -f /etc/lvmtab ]
then
    action $"Setting up Logical Volume Management:" \
        /sbin/vgscan && /sbin/vgchange -ay
fi
```

## 2. Edit the file `/etc/rc.d/init.d/halt` to deactivate the volume groups at shutdown.

Insert the `vgchange -an` command near the end of this.

---

# LVM Files

---

- `/etc/lvmtab`

Informs about the activated VGs

- `/proc/lvm/*`

Inform about the VGs structure

- `/dev/my_VG[1-n]/*`

LVM device files

---

# LVM Commands 1/7

---

- **Initializing** disks or disk partitions

- For entire disks: `# pvcreate /dev/hdb`

- For partitions: set the partition type to 0x8e using fdisk. Next:

- `# pvcreate /dev/hdb1`

- `pvcreate – physical volume "/dev/hdb1" successfully created`

- **Creating a VG:**

- `# vgcreate my_VG /dev/hdb1 [/dev/hdc1] ...`

- `vgcreate – INFO: using default physical extent size 4 MB (default)`

- `vgcreate – INFO: maximum logical volume size is 255.99 Gigabyte`

- `vgcreate – doing automatic backup of volume group "my_VG"`

- `vgcreate – volume group "my_VG" successfully created and activated`

- `# more /etc/lvmtab`

- `my_VG`

---

# LVM Commands 2/7

---

- **Activating a VG:** `# vgchange -ay my_VG`

**Note:** in the startup script

- **Removing a VG:**

1. Deactivate the VG: `# vgchange -an my_VG`

2. Remove the VG: `# vgremove my_VG`

- **Adding PVs to a VG:**

`# pvcreate /dev/hdb2`

`# vgextend my_VG /dev/hdb2`

- **Removing PVs from a VG:**

`# vgreduce my_VG /dev/hdb2`

---

# LVM Commands 3/7

---

- **Creating an LV** in `my_VG`: `# lvcreate options -nmy_LV`

Main options:

`-c y/n` Linear. Default is no Linear (Striped)

`-i PV_Number` number of PVs to scatter the LV.

`-I S_Size` Strip size (in KBytes). The strip is the transactions unit.

`-l LEs_Number` number of LEs (LE size = PE size) for the new LV.

`-L LV_Size[kKmMgGtT]` size for the new LV. K (kilobytes), M (megabytes), G (gigabytes) or T (terabytes). Default unit is megabytes.

`-n LV_Name`

`-p r/w` permission

---

# LVM Commands 4/7

---

- To create a 1500MB linear LV named `my_LV` in `my_VG`:

```
# lvcreate -Cy -L1500 -nmy_LV my_VG
```

- To create a LV of size 100 LEs (or PEs), with 2 PVs and strip size of 4 KB:

```
# lvcreate -i2 -I4 -l100 -nmy_LV my_VG
```

```
# vgdisplay my_VG
```

— Volume group —

VG Name my\_VG

VG Access read/write

VG Status available/resizable

VG # 0

MAX LV 256

-----

---

# LVM Commands 5/7

---

VG Size 1.95 GB

PE Size 4 MB

Total PE 498

Alloc PE / Size 100 / 400 MB

Free PE / Size 398 / 1.55 GB

- If you want to create an LV that uses the entire VG, use `vgdisplay` to find the "Total PE" size, then use that when running `lvcreate`:

```
# vgdisplay my_VG | grep "Total PE"
```

Total PE 10230

```
# lvcreate -l10230 my_VG -nmy_LV
```

# LVM Commands 6/7

---

- **Create an ext2** file system on the LV

```
# mke2fs /dev/my_VG/my_LV
```

- **Mount and Test** the File System

```
# mount -t auto /dev/my_VG/my_LV /mnt/LV1
```

```
# df
```

| Filesystem       | 1k-blocks | Used     | Available | Use% | Mounted on |
|------------------|-----------|----------|-----------|------|------------|
| /dev/hda1        | 35886784  | 13360620 | 20703192  | 40%  | /          |
| /dev/my_VG/my_LV | 396672    | 13       | 376179    | 1%   | /mnt/LV1   |

- Obtaining **LV** information

```
# lvdisplay /dev/my_VG/my_LV (strip size 4KB)
```

- **Removing a LV**

```
# umount /dev/my_VG/my_LV && lvremove /dev/my_VG/my_LV
```



- **Extending (+4MBs) a LV** (/dev/my\_VG/my\_LV) mounted on /mnt/LV1 :

```
# lvextend -L+4M /dev/my_VG/my_LV
```

```
# umount /dev/my_VG/my_LV
```

```
# e2fsck -f /dev/my_VG/my_LV
```

```
# resize2fs /dev/my_VG/my_LV
```

```
# mount /dev/my_VG/my_LV /mnt/LV1
```

- **Reducing (-4MBs) a LV** (/dev/my\_VG/my\_LV) mounted on /mnt/LV1 :

```
# lvreduce -L-4M /dev/my_VG/my_LV
```

```
# umount /mnt/LV1
```

```
# e2fsck -f /dev/my_VG/my_LV
```

```
# resize2fs /dev/my_VG/my_LV
```

```
# mount /dev/my_VG/my_LV /mnt/LV1
```

---

# Examples

---

- **Setting up LVM on three SCSI disks (/dev/sda, /dev/sdb, and /dev/sdc) with striping**
- **Taking a Backup Using Snapshots**
- **Removing an Old Disk (/dev/hdb)**
- **Distributing Old Extents to a New Replacement Disk**

---

# Setting up LVM on three SCSI disks (/dev/sda, /dev/sdb, and /dev/sdc) with striping 1/2

---

- Preparing the disk partitions

```
# pvcreate /dev/sda
```

```
# pvcreate /dev/sdb
```

```
# pvcreate /dev/sdc
```

- Setup a VG

1. Create a VG

```
# vgcreate my_VG /dev/sda /dev/sdb /dev/sdc
```

2. Run `vgdisplay` to verify volume group

# Setting up LVM on three SCSI disks (/dev/sda, /dev/sdb, and /dev/sdc) with striping 2/2

---

- **Creating** the **LV** (1GB) with 3 PVs and strip size of 4 KB on the VG:

```
# lvcreate -i3 -I4 -L1G -nmy_LV my_VG
```

- **Create** an **ext2** file system on the LV

```
# mke2fs /dev/my_VG/my_LV
```

- **Mount** and **Test** the File System

```
# mount /dev/my_VG/my_LV /mnt
```

```
# df
```

---

# Taking a Backup Using Snapshots

---

1. **Create** the snapshot volume of 500MB (free space in **my\_VG**):

```
# lvcreate -L500M -s -nmy_LV_backup /dev/my_VG/my_LV
```

2. **Mount** the snapshot volume:

```
# mkdir /mnt/my_LV_backup
```

```
# mount /dev/my_VG/my_LV_backup /mnt/my_LV_backup
```

3. **Remove** the snapshot:

```
# umount /mnt/my_LV_backup
```

```
# lvremove /dev/my_VG/my_LV_backup
```

---

# Removing an Old Disk (/dev/hdb)

---

## 1. Distributing Old Extents to Existing Disks in VG

```
# pvmove /dev/hdb
```

## 2. Remove the unused disk

```
# vgreduce my_VG /dev/hdb
```

## 3. The drive can now be removed when the machine is powered down

# Distributing Old Extents to a New Replacement Disk

---

## 1. Prepare the new disk

```
# pvcreate /dev/sdf
```

## 2. Add it to the VG `my_VG`

```
# vgextend my_VG /dev/sdf
```

## 3. Move the data

```
# pvmove /dev/hdb /dev/sdf
```

## 4. Remove the unused disk

```
# vgreduce my_VG /dev/hdb
```

## 5. The drive can now be removed when the machine is powered down

# Kernel configuration

